

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error
```

```
In [2]: data = pd.read_csv('http://bit.ly/w-data')
data.head(10)
```


```
Out[2]:
```

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	1.5	20
6	9.2	88
7	5.5	60
8	8.3	81
9	2.7	25


```
In [3]: data.isnull == True
```

```
Out[3]: False
```

```
In [4]: sns.set_style('darkgrid')
sns.scatterplot(y= data['Scores'], x= data['Hours'])
plt.title('Marks Vs Study Hours',size=20)
plt.ylabel('Marks Percentage', size=12)
plt.xlabel('Hours Studied', size=12)
plt.show()
```



```
In [5]: sns.regplot(x= data['Hours'], y= data['Scores'])
plt.title('Regression Plot',size=20)
plt.ylabel('Marks Percentage', size=12)
plt.xlabel('Hours Studied', size=12)
plt.show()
print(data.corr())
```



	Hours	Scores
Hours	1.000000	0.976191
Scores	0.976191	1.000000

```
In [6]: X = data.iloc[:, :-1].values
        y = data.iloc[:, 1].values

        # Splitting the Data in two
        train_X, val_X, train_y, val_y = train_test_split(X, y, random_state = 0)
```

```
In [7]: regression = LinearRegression()
        regression.fit(train_X, train_y)
        print("-----Model Trained-----")

        -----Model Trained-----
```

```
In [8]: pred_y = regression.predict(val_X)
        prediction = pd.DataFrame({'Hours': [i[0] for i in val_X], 'Predicted Marks': [k for k in pred_y]})
        prediction
```

```
Out[8]:
```


	Hours	Predicted Marks
0	1.5	16.844722
1	3.2	33.745575
2	7.4	75.500624
3	2.5	26.786400
4	5.9	60.588106
5	3.8	39.710582
6	1.9	20.821393

```
In [9]: compare_scores = pd.DataFrame({'Actual Marks': val_y, 'Predicted Marks': pred_y})
        compare_scores
```

```
Out[9]:
```

	Actual Marks	Predicted Marks
0	20	16.844722
1	27	33.745575
2	69	75.500624
3	30	26.786400
4	62	60.588106
5	35	39.710582
6	24	20.821393

```
In [10]: plt.scatter(x=val_X, y=val_y, color='blue')
        plt.plot(val_X, pred_y, color='Black')
        plt.title('Actual vs Predicted', size=20)
        plt.ylabel('Marks Percentage', size=12)
        plt.xlabel('Hours Studied', size=12)
        plt.show()
```



```
In [11]: print('Mean absolute error: ', mean_absolute_error(val_y, pred_y))

Mean absolute error: 4.130879918502486
```

```
In [12]: hours = [9.25]
        answer = regression.predict([hours])
        print("Score = {}".format(round(answer[0],3)))

Score = 93.893
```

In [ ]:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js